CalibrateEmulateSample.jl: Accelerated Parametric Uncertainty Quantification

Oliver R. A. Dunbar [1,¶], **Melanie Bieli**[2], **Alfredo Garbuno-Iñigo** [3], **Michael Howland** [4], **Andre Nogueira de Souza** [5], **Laura Anne Mansfield** [6], **Gregory L. Wagner** [5], **and N. Efrat-Henrici**[1]

1 Geological and Planetary Sciences, California Institute of Technology 2 Swiss Re Ltd. 3 Department of Statistics, Mexico Autonomous Institute of Technology 4 Civil and Environmental Engineering, Massachusetts Institute of Technology 5 Earth, Atmospheric, and Planetary Sciences, Massachusetts Institute of Technology 6 Earth System Science, Doerr School of Sustainability, Stanford University ¶ Corresponding author

Summary

A Julia language (Bezanson et al., 2017) package providing practical and modular implementation of "Calibrate, Emulate, Sample" (Cleary et al., 2021), hereafter CES, an accelerated workflow for obtaining model parametric uncertainty is presented. This is also known as Bayesian inversion or uncertainty quantification. To apply CES one requires a computer model (written in any programming language) dependent on free parameters, a prior distribution encoding some prior knowledge about the distribution over the free parameters, and some data with which to constrain this prior distribution. The pipeline has three stages, most easily explained in reverse: 1. the final stage, and goal of the workflow, is to draw samples (Sample) from the Bayesian posterior distribution, that is, the prior distribution conditioned on the observed data; 2. to accelerate and regularize this process we train statistical emulators to represent the user-provided parameter-to-data map (Emulate); 3. the training points for these emulators are generated by the computer model, and selected adaptively around regions of high posterior mass (Calibrate). We describe CES as an accelerated workflow, as it is often able to use dramatically fewer evaluations of the computer model when compared with applying sampling algorithms, such as Markov chain Monte Carlo (MCMC), directly.

- Calibration tools: We recommend choosing adaptive training points with Ensemble Kalman methods such as EKI (Iglesias et al., 2013) and its variants (Huang et al., 2022); and CES provides explicit utilities from the codebase EnsembleKalmanProcesses.jl (Dunbar, Lopez-Gomez, et al., 2022).
- Emulation tools: CES integrates any statistical emulator, currently implemented are Gaussian Processes (GP) (Williams & Rasmussen, 2006), explicitly provided through packages SciKitLearn.jl (Pedregosa et al., 2011) and GaussianProcesses.jl (Fairbrother et al., 2022), and Random Features (Liu et al., 2022; Rahimi et al., 2007; Rahimi & Recht, 2008), explicitly provided through RandomFeatures.jl that can provide additional flexibility and scalability, particularly in higher dimensions.
- Sampling tools: The regularized and accelerated sampling problem is solved with MCMC, and CES provides the variants of Random Walk Metropolis (Metropolis et al., 1953; Sherlock et al., 2010), and preconditioned Crank-Nicholson (Cotter et al., 2013), using APIs from Turing.jl. Some regular emulator mean functions are differentiable, and including accelerations of derivative-based MCMC into CES, (e.g., NUTS, Hoffman et al., 2014; Barker, Livingstone & Zanella, 2022); is an active direction of work.

To highlight code accessibility, we also provide a suite of detailed scientifically-inspired examples, with documentation that walks users through some use cases. Such use cases not only

45 demonstrate the capability of the CES pipeline, but also teach users about typical interface
46 and workflow experience.

Statement of need

48 Computationally expensive computer codes for predictive modelling are ubiquitous across
49 science and engineering disciplines. Free parameter values that exist within these modelling
50 frameworks are typically constrained by observations to produce accurate and robust predictions
51 about the system they are approximating numerically. In a Bayesian setting, this is viewed
52 as evolving an initial parameter distribution (based on prior information) with the input of
53 observed data, to a more informative data-consistent distribution (posterior). Unfortunately,
54 this task is intensely computationally expensive, commonly requiring over 10^5 evaluations of
55 the expensive computer code (e.g., Random Walk Metropolis), with accelerations relying on
56 intrusive model information, such as a derivative of the parameter-to-data map. CES is able
57 to approximate and accelerate this process in a non-intrusive fashion and requiring only on the
58 order of 10^2 evaluations of the original computer model. This opens the doors for quantifying
59 parametric uncertainty for a class of numerically intensive computer codes that has previously
60 been unavailable.

State of the field

62 In Julia there are a few tools for performing non-accelerated uncertainty quantification, from
63 classical sensitivity analysis approaches, for example, UncertaintyQuantification.jl, GlobalSen-
64 sitivity.jl (Dixit & Rackauckas, 2022), and MCMC, for example, Mamba.jl or Turing.jl. For
65 computational efficiency, ensemble methods also provide approximate sampling, (Dunbar,
66 Lopez-Gomez, et al., 2022; e.g., the Ensemble Kalman Sampler Garbuno-Inigo et al., 2020),
67 though these only provide Gaussian approximations of the posterior.

68 Accelerated uncertainty quantification tools also exist for the related approach of Approximate
69 Bayesian Computation (ABC), for example, GpABC (Tankhilevich et al., 2020) or Approx-
70 Bayes.jl; these tools both approximately sample from the posterior distribution. In ABC, this
71 approximation comes from bypassing the likelihood that is usually required in sampling methods,
72 such as MCMC. Instead, the goal of ABC is to replace the likelihood with a scalar-valued
73 sampling objective that compares model and data. In CES, the approximation comes from
74 learning the parameter-to-data map, then following this it calculates an explicit likelihood and
75 uses exact sampling via MCMC. Some ABC algorithms also make use of statistical emulators
76 to further accelerate sampling (GpABC). Although flexible, ABC encounters challenges due to
77 the subjectivity of summary statistics and distance metrics, that may lead to approximation
78 errors particularly in high-dimensional settings (Nott et al., 2018). CES is more restrictive due
79 to use of an explicit Gaussian likelihood, but also leverages this structure to deal with high
80 dimensional data.

81 Several other tools are available in other languages for a purpose of accelerated learning of the
82 posterior distribution or posterior sampling. Two such examples, written in Python, approximate
83 the log-posterior distribution directly with a Gaussian process: PyVBMC (Huggins et al., 2023)
84 additionaly uses variational approximations to calculate the normalization constant, and GPry
85 (Gammal et al., 2023), which iteratively trains the GP with an active training point selection
86 algorithm. Such algorithms are distinct from CES, which approximates the parameter-to-data
87 map with the Gaussian process, and advocates ensemble Kalman methods to select training
88 points.

A simple example from the code documentation

We sketch an end-to-end example of the pipeline, with fully-detailed walkthrough given in the online documentation.

We have a model of a sinusoidal signal that is a function of parameters $\theta = (A, v)$, where A is the amplitude of the signal and v is vertical shift of the signal

$$f(A, v) = A\sin(\phi + t) + v, \forall t \in [0, 2\pi].$$

Here, ϕ is the random phase of each signal. The goal is to estimate not just point estimates of the parameters $\theta = (A, v)$, but entire probability distributions of them, given some noisy observations. We will use the range and mean of a signal as our observable:

$$G(\theta) = [\text{range}(f(\theta)), \text{mean}(f(\theta))]$$

Then, our noisy observations, y_{obs}, can be written as:

$$y_{obs} = G(\theta^\dagger) + \mathcal{N}(0, \Gamma)$$

where Γ is the observational covariance matrix. We will assume the noise to be independent for each observable, giving us a diagonal covariance matrix.


Figure 1: The true and observed range and mean.

For this experiment $\theta^\dagger = (A^\dagger, v^\dagger) = (3.0, 7.0)$, and the noisy observations are displayed in blue in Figure 1.

We define prior distributions on the two parameters. For the amplitude, we define a prior with mean 2 and standard deviation 1. It is additionally constrained to be nonnegative. For the vertical shift we define a prior with mean 0 and standard deviation 5.

```
const PD = CalibrateEmulateSample.ParameterDistributions
prior_u1 = PD.constrained_gaussian("amplitude", 2, 1, 0, Inf)
prior_u2 = PD.constrained_gaussian("vert_shift", 0, 5, -Inf, Inf)
prior = PD.combine_distributions([prior_u1, prior_u2])
```

Figure 2: Marginal distributions of the prior

¹⁰⁵ The prior is displayed in Figure 2.

¹⁰⁶ We now adaptively find input-output pairs from our map G in a region of interest using an
¹⁰⁷ inversion method (an ensemble Kalman process). This is the Calibrate stage, and iteratively
¹⁰⁸ generates parameter combinations, that refine around a region of high posterior mass.

```julia
const EKP = CalibrateEmulateSample.EnsembleKalmanProcesses
N_ensemble = 10
N_iterations = 5
initial_ensemble = EKP.construct_initial_ensemble(prior, N_ensemble)
ensemble_kalman_process = EKP.EnsembleKalmanProcess(
    initial_ensemble, y_obs, Γ, EKP.Inversion();
)
for i in 1:N_iterations
    params_i = EKP.get_phi_final(prior, ensemble_kalman_process)
    G_ens = hcat([G(params_i[:, i]) for i in 1:N_ensemble]...)
    EKP.update_ensemble!(ensemble_kalman_process, G_ens)
end
```


Figure 3: The resulting ensemble from a calibration.

¹⁰⁹ The adaptively refined training points from EKP are displayed in Figure 3. We now build an
¹¹⁰ basic Gaussian process emulator from the GaussianProcesses.jl package to emulate the map G
¹¹¹ using these points.

```julia
const UT = CalibrateEmulateSample.Utilities
const EM = CalibrateEmulateSample.Emulators
```

```
input_output_pairs = UT.get_training_points(
    ensemble_kalman_process, N_iterations,
)
gppackage = EM.GPJL()
gauss_proc = EM.GaussianProcess(gppackage, noise_learn = false)
emulator = EM.Emulator(
    gauss_proc, input_output_pairs, normalize_inputs = true,  obs_noise_cov = Γ,
)
EM.optimize_hyperparameters!(emulator) # train the emulator
```


Figure 4: The Gaussian process emulator of the range and mean maps, trained on the re-used calibration pairs

We evaluate the mean of this emulator on a grid, and also show the value of the true G at training point locations in Figure 4.

We can then sample with this emulator using an MCMC scheme. We first choose a good step size (an algorithm parameter) by running some short sampling runs (of length 2,000 steps). Then we run the 100,000 step sampling run to generate samples of the joint posterior distribution.

```
const MC = CalibrateEmulateSample.MarkovChainMonteCarlo
mcmc = MC.MCMCWrapper(
    MC.RWMHSampling(), y_obs, prior, emulator,
)
# choose a step size
new_step = MC.optimize_stepsize(
    mcmc; init_stepsize = 0.1, N = 2000,
)
# Now begin the actual MCMC
chain = MC.sample(
    mcmc, 100_000; stepsize = new_step, discard_initial = 2_000,
)
```

Figure 5: The joint posterior distribution histogram

A histogram of the samples from the CES algorithm is displayed in Figure 5. We see that the posterior distribution contains the true value $(3.0, 7.0)$ with high probability.

Research projects using the package

Some research projects that use this codebase, or modifications of it, are * (Dunbar et al., 2021) * (Bieli et al., 2022) * (Hillier, 2022) * (Howland et al., 2022) * (Dunbar, Howland, et al., 2022) * (Mansfield & Sheshadri, 2022) * (King et al., 2023)

Acknowledgements

References

Bezanson, J., Edelman, A., Karpinski, S., & Shah, V. B. (2017). Julia: A fresh approach to numerical computing. *SIAM Review*, *59*(1), 65–98. https://doi.org/10.1137/141000671

Bieli, M., Dunbar, O. R. A., Jong, E. K. de, Jaruga, A., Schneider, T., & Bischoff, T. (2022). An efficient Bayesian approach to learning droplet collision kernels: Proof of concept using "Cloudy," a new n-moment bulk microphysics scheme. *Journal of Advances in Modeling Earth Systems*, *14*(8), e2022MS002994. https://doi.org/10.1029/2022MS002994

Cleary, E., Garbuno-Inigo, A., Lan, S., Schneider, T., & Stuart, A. M. (2021). Calibrate, emulate, sample. *Journal of Computational Physics*, *424*, 109716. https://doi.org/10.1016/j.jcp.2020.109716

Cotter, S. L., Roberts, G. O., Stuart, A. M., & White, D. (2013). MCMC Methods for Functions: Modifying Old Algorithms to Make Them Faster. *Statistical Science*, *28*(3), 424–446. https://doi.org/10.1214/13-STS421

Dixit, V. K., & Rackauckas, C. (2022). GlobalSensitivity.jl: Performant and parallel global sensitivity analysis with julia. *Journal of Open Source Software*, *7*(76), 4561. https://doi.org/10.21105/joss.04561

Dunbar, O. R. A., Garbuno-Inigo, A., Schneider, T., & Stuart, A. M. (2021). Calibration and uncertainty quantification of convective parameters in an idealized GCM. *Journal of Advances in Modeling Earth Systems*, *13*(9), e2020MS002454. https://doi.org/10.1029/2020MS002454

Dunbar, O. R. A., Howland, M. F., Schneider, T., & Stuart, A. M. (2022). Ensemble-based experimental design for targeting data acquisition to inform climate models. *Journal of Advances in Modeling Earth Systems*, *14*(9), e2022MS002997. https://doi.org/10.1029/2022MS002997

Dunbar, O. R. A., Lopez-Gomez, I., Garbuno-Iñigo, A. G.-I., Huang, D. Z., Bach, E., & Wu, J. (2022). EnsembleKalmanProcesses.jl: Derivative-free ensemble-based model calibration. *Journal of Open Source Software*, *7*(80), 4869. https://doi.org/10.21105/joss.04869

Fairbrother, J., Nemeth, C., Rischard, M., Brea, J., & Pinder, T. (2022). GaussianProcesses. Jl: A nonparametric bayes package for the julia language. *Journal of Statistical Software*, *102*, 1–36. https://doi.org/10.18637/jss.v102.i01

Gammal, J. E., Schöneberg, N., Torrado, J., & Fidler, C. (2023). Fast and robust bayesian inference using gaussian processes with GPry. *Journal of Cosmology and Astroparticle Physics*, *2023*(10), 021. https://doi.org/10.1088/1475-7516/2023/10/021

Garbuno-Inigo, A., Nüsken, N., & Reich, S. (2020). Affine invariant interacting Langevin dynamics for Bayesian inference. *SIAM Journal on Applied Dynamical Systems*, *19*(3), 1633–1658. https://doi.org/10.1137/19M1304891

Hillier, A. (2022). *Supervised calibration and uncertainty quantification of subgrid closure parameters using ensemble Kalman inversion* [Master's thesis, Massachusetts Institute of Technology. Department of Electrical Engineering; Computer Science]. https://doi.org/1721.1/145140

Hoffman, M. D., Gelman, A., & others. (2014). The no-u-turn sampler: Adaptively setting path lengths in hamiltonian monte carlo. *J. Mach. Learn. Res.*, *15*(1), 1593–1623.

Howland, M. F., Dunbar, O. R. A., & Schneider, T. (2022). Parameter uncertainty quantification in an idealized GCM with a seasonal cycle. *Journal of Advances in Modeling Earth Systems*, *14*(3), e2021MS002735. https://doi.org/10.1029/2021MS002735

Huang, D. Z., Huang, J., Reich, S., & Stuart, A. M. (2022). Efficient derivative-free bayesian inference for large-scale inverse problems. *Inverse Problems*, *38*(12), 125006. https://doi.org/10.1088/1361-6420/ac99fa

Huggins, B., Li, C., Tobaben, M., Aarnos, M. J., & Acerbi, L. (2023). PyVBMC: Efficient bayesian inference in python. *Journal of Open Source Software*, *8*(86), 5428. https://doi.org/10.21105/joss.05428

Iglesias, M. A., Law, K. J., & Stuart, A. M. (2013). Ensemble kalman methods for inverse problems. *Inverse Problems*, *29*(4), 045001. https://doi.org/10.1088/0266-5611/29/4/045001

King, R. C., Mansfield, L. A., & Sheshadri, A. (2023). *Bayesian history matching applied to the calibration of a gravity wave parameterization* [Preprint]. https://doi.org/10.22541/essoar.170365299.96491153/v1

188 Liu, F., Huang, X., Chen, Y., & Suykens, J. A. K. (2022). Random features for kernel
189 approximation: A survey on algorithms, theory, and beyond. *IEEE Transactions on Pattern*
190 *Analysis and Machine Intelligence*, *44*(10), 7128–7148. https://doi.org/10.1109/TPAMI.
191 2021.3097011

192 Livingstone, S., & Zanella, G. (2022). The barker proposal: Combining robustness and
193 efficiency in gradient-based MCMC. *Journal of the Royal Statistical Society Series B:*
194 *Statistical Methodology*, *84*(2), 496–523. https://doi.org/10.1111/rssb.12482

195 Mansfield, L. A., & Sheshadri, A. (2022). Calibration and uncertainty quantification of
196 a gravity wave parameterization: A case study of the Quasi-Biennial Oscillation in an
197 intermediate complexity climate model. *Journal of Advances in Modeling Earth Systems*,
198 *14*(11). https://doi.org/10.1029/2022MS003245

199 Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., & Teller, E. (1953).
200 Equation of state calculations by fast computing machines. *The Journal of Chemical*
201 *Physics*, *21*(6), 1087–1092. https://doi.org/10.1063/1.1699114

202 Nott, D. J., Ong, V. M.-H., Fan, Y., & Sisson, S. A. (2018). High-Dimensional ABC. In
203 *Handbook of Approximate Bayesian Computation* (pp. 211–241). CRC Press. https:
204 //doi.org/10.1201/9781315117195-8

205 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M.,
206 Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D.,
207 Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine learning in Python.
208 *Journal of Machine Learning Research*, *12*, 2825–2830.

209 Rahimi, A., & Recht, B. (2008). Uniform approximation of functions with random bases. *2008*
210 *46th Annual Allerton Conference on Communication, Control, and Computing*, 555–561.
211 https://doi.org/10.1109/allerton.2008.4797607

212 Rahimi, A., Recht, B., & others. (2007). Random features for large-scale kernel
213 machines. *NIPS*, *3*, 5. https://proceedings.neurips.cc/paper_files/paper/2007/file/
214 013a006f03dbc5392effeb8f18fda755-Paper.pdf

215 Sherlock, C., Fearnhead, P., & Roberts, G. O. (2010). The random walk metropolis: Linking
216 theory and practice through a case study. *Statistical Science*, *25*(2), 172–190. https:
217 //doi.org/10.1214/10-sts327

218 Tankhilevich, E., Ish-Horowicz, J., Hameed, T., Roesch, E., Kleijn, I., Stumpf, M. P. H., & He,
219 F. (2020). GpABC: a Julia package for approximate Bayesian computation with Gaussian
220 process emulation. *Bioinformatics*. https://doi.org/10.1093/bioinformatics/btaa078

221 Williams, C. K., & Rasmussen, C. E. (2006). *Gaussian processes for machine learning* (Vol.
222 2). MIT press Cambridge, MA. https://doi.org/10.1142/S0129065704001899